

+



CHULA ENGINEERING
Foundation toward Innovation

COMPUTER

Chula Big Data and IoT
Center of Excellence
(CUBIC)



Data Analytics (Part2) Python for Data Analytics

Peerapon Vateekul, Ph.D.

Department of Computer Engineering,
Faculty of Engineering, Chulalongkorn University
Peerapon.v@chula.ac.th



Outlines

- Lab (cont.)
- Collaborative Filtering

Python for Data Science and Machine Learning Bootcamp

Jose Marcial Portilla

\$120

PIERIAN DATA

<https://www.pieriandata.com/>

- Understand data analytics **tasks**
- Be able to identify tasks and **tools** (technique) from a given problem

+

Lab (cont.)



Lab9: Association Rule

Online Retail Data Set



- A useful (but somewhat overlooked) technique is called association analysis which attempts to find common patterns of items in large data sets. One specific application is often called market basket analysis.

- 541909 rows and 8 columns

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	2010-12-01 08:26:00	2.55	17850.0	United Kingdom
1	536365	71053	WHITE METAL LANTERN	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	2010-12-01 08:26:00	2.75	17850.0	United Kingdom
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom

■ Library

```
# Build up the frequent items
frequent_itemsets = apriori(basket_sets, min_support=0.07, use_colnames=True)
```

```
# Create the rules
rules = association_rules(frequent_itemsets, metric="lift", min_threshold=1)
rules
```

■ Output image

	antecedants	consequents	antecedent support	consequent support	support	confidence	lift	leverage	conviction
2	(ALARM CLOCK BAKELIKE GREEN)	(ALARM CLOCK BAKELIKE RED)	0.096939	0.094388	0.079082	0.815789	8.642959	0.069932	4.916181
3	(ALARM CLOCK BAKELIKE RED)	(ALARM CLOCK BAKELIKE GREEN)	0.094388	0.096939	0.079082	0.837838	8.642959	0.069932	5.568878
17	(SET/6 RED SPOTTY PAPER PLATES)	(SET/20 RED RETROSPOT PAPER NAPKINS)	0.127551	0.132653	0.102041	0.800000	6.030769	0.085121	4.336735
18	(SET/6 RED SPOTTY PAPER CUPS)	(SET/6 RED SPOTTY PAPER PLATES)	0.137755	0.127551	0.122449	0.888889	6.968889	0.104878	7.852041

Reference: <http://pbpython.com/market-basket-analysis.html>

+ Lab10: Time Series Forecasting with ARIMA: CO2 Data Set



- We'll be working with a dataset called "**Atmospheric CO2** from Continuous Air Samples at Mauna Loa Observatory, Hawaii, U.S.A.," which collected CO2 samples from March 1958 to December 2001. 1969 rows and 2 columns

co2	
1958-03-29	316.1
1958-04-05	317.3
1958-04-12	317.6
1958-04-19	317.5
1958-04-26	316.4
1958-03-01	316.100000
1958-04-01	317.200000
1958-05-01	317.433333
1958-06-01	315.625000
1958-07-01	315.625000
Freq: MS, Name: co2, dtype: float64	

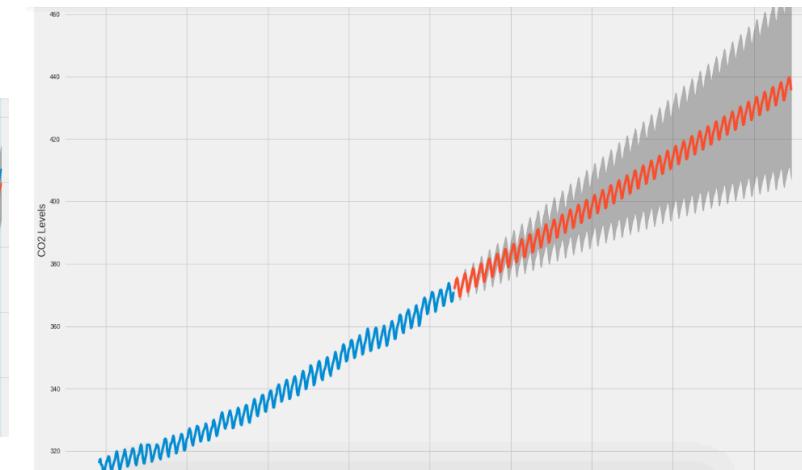
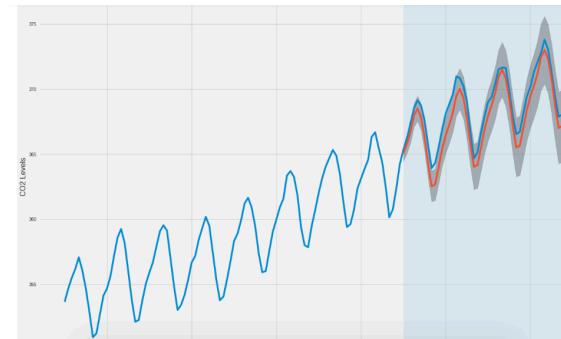
■ Library

Fitting an ARIMA Time Series Model

```
mod = sm.tsa.statespace.SARIMAX(y,
                                 order=(1, 1, 1),
                                 seasonal_order=(1, 1, 1, 12),
                                 enforce_stationarity=False, # Whether or not
                                 enforce_invertibility=False) # Whether or not
```

results = mod.fit()
print(results.summary())

■ Output image





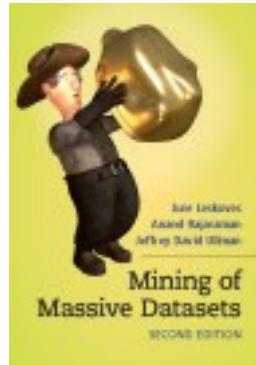
Recommender System: Collaborative Filtering



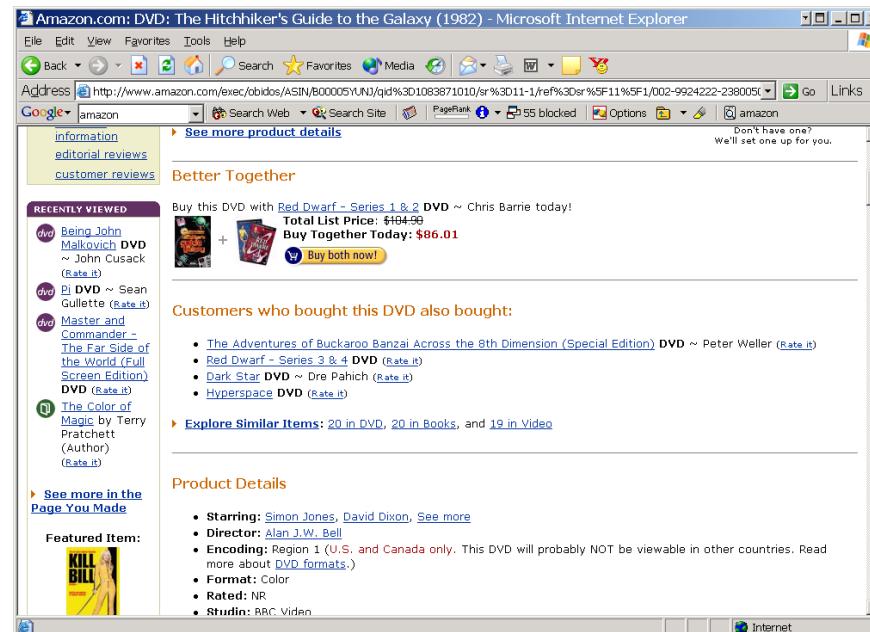
Acknowledgement

- Jure Leskovec, Anand Rajaraman, Jeff Ullman
- Stanford University
- Mining of Massive Datasets
- www.mmds.org

- Prof. Dr. Boonserm Kijsirikul
- 2110746 Big Data Analytics

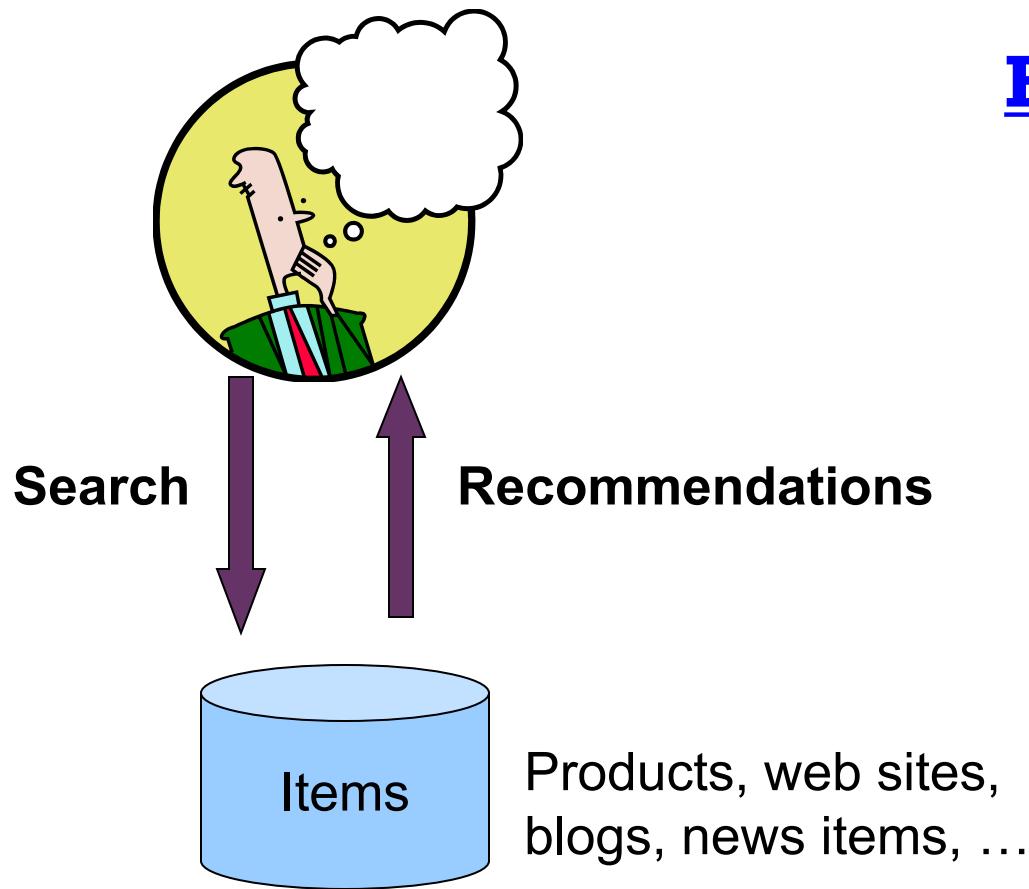


■Introduction to Recommender System





Recommendations



Examples:

amazon.com.



StumbleUpon



movie lens
helping you find the *right* movies

last.fm™
the social music revolution



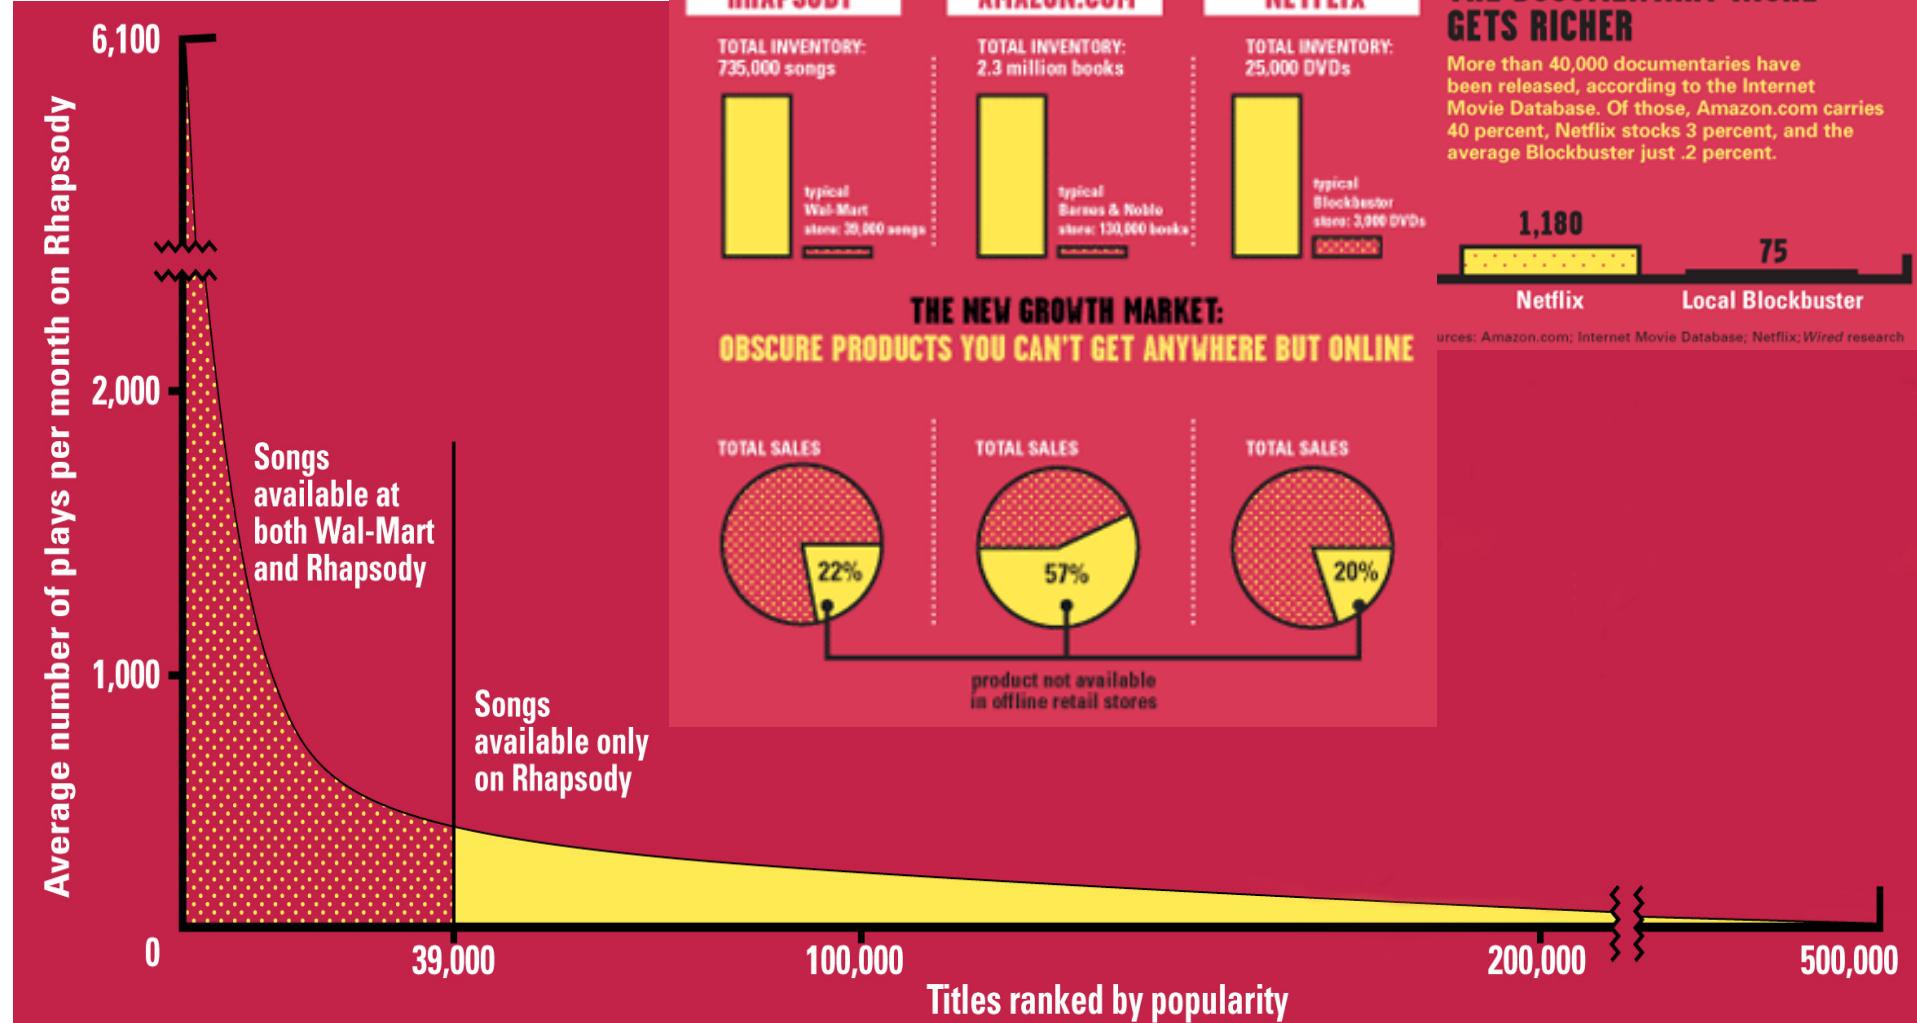


From Scarcity to Abundance

- **Shelf space is a scarce commodity for traditional retailers**
 - Also: TV networks, movie theaters,...
- **Web enables near-zero-cost dissemination of information about products**
 - From scarcity to abundance
- **More choice necessitates better filters**
 - Recommendation engines
 - How **Into Thin Air** made **Touching the Void** a bestseller: <http://www.wired.com/wired/archive/12.10/tail.html>



Sidenote: The Long Tail



Source: Chris Anderson (2004), MIT, and Michael Smith, Carnegie Mellon; Barnes & Noble; Netflix; RealNetworks

Types of Recommendations

■ Editorial and hand curated

- List of favorites
- Lists of “essential” items

■ Simple aggregates

- Top 10, Most Popular, Recent Uploads

■ Tailored to individual users

- Amazon, Netflix, ...



How to recommend to each user?

- 1) Content-based
- **2) Collaborative**
- 3) Latent factor based

Today!



1. User-based CF
2. Item-based CF

■ Collaborative Filtering (CF)

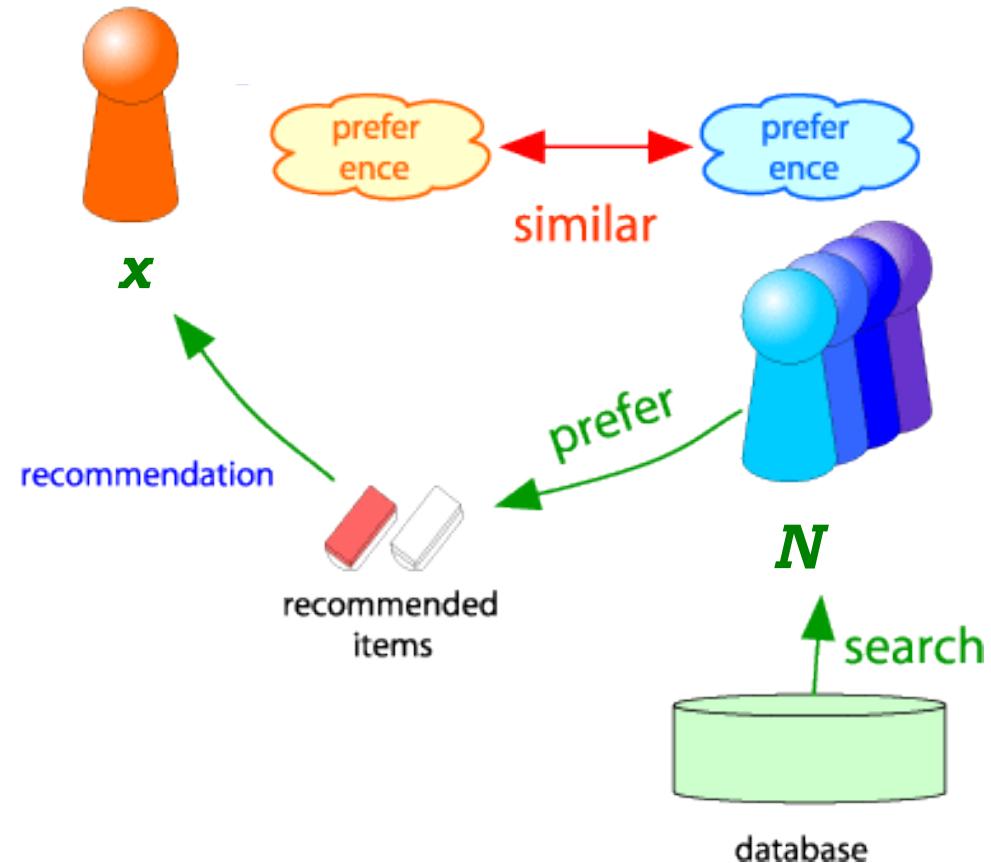
1	3	4		
	3	5		5
		4	5	5
			3	
1				





User-Based CF

- Consider user x
- Find set N of other users whose ratings are “**similar**” to x ’s ratings
- Estimate x ’s ratings based on ratings of users in N



+ Finding “Similar” Users

- Let r_x be the vector of user x 's ratings
- Jaccard similarity measure**
 - Problem:** Ignores the value of the rating
- Cosine similarity measure**
 - $\text{sim}(x, y) = \cos(r_x, r_y) = \frac{r_x \cdot r_y}{\|r_x\| \cdot \|r_y\|}$
 - Problem:** Treats missing ratings as “negative”

$$\begin{aligned}r_x &= [* , _, _, *, _, ***] \\r_y &= [* , _, **, **, _,]\end{aligned}$$

r_x, r_y as sets:
 $r_x = \{1, 4, 5\}$
 $r_y = \{1, 3, 4\}$

r_x, r_y as points:
 $r_x = \{1, 0, 0, 1, 3\}$
 $r_y = \{1, 0, 2, 2, 0\}$



Similarity Metric

Cosine sim:

$$\text{sim}(x, y) = \frac{\sum_i r_{xi} \cdot r_{yi}}{\sqrt{\sum_i r_{xi}^2} \cdot \sqrt{\sum_i r_{yi}^2}}$$

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

- **Intuitively we want:** $\text{sim}(A, B) > \text{sim}(A, C)$

- **Jaccard similarity:** $1/5 < 2/4$

- **Cosine similarity:** $0.386 > 0.322$

- Considers missing ratings as “negative”
- **Solution: subtract the (row) mean**

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	2/3			5/3	-7/3		
B	1/3	1/3	-2/3				
C				-5/3	1/3	4/3	
D		0					0

sim A,B vs. A,C:
0.092 > -0.559

Notice cosine sim. is correlation when data is centered at 0



Rating Predictions

From similarity metric to recommendations:

- Let r_x be the vector of user x 's ratings
- Let N be the set of k users most similar to x who have rated item i
- **Prediction for item s of user x :**
 - $r_{xi} = \frac{1}{k} \sum_{y \in N} r_{yi}$
 - $r_{xi} = \frac{\sum_{y \in N} s_{xy} \cdot r_{yi}}{\sum_{y \in N} s_{xy}}$
 - Other options?
- **Many other tricks possible...**

Shorthand:

$$s_{xy} = sim(x, y)$$



CF ($|N|=2$): Train

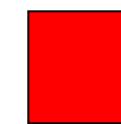
	users											
	1	2	3	4	5	6	7	8	9	10	11	12
1	1		3			5			5		4	
2			5	4			4			2	1	3
3	2	4		1	2		3		4	3	5	
4		2	4		5			4			2	
5			4	3	4	2					2	5
6	1		3		3			2			4	

 - unknown rating
 - rating between 1 to 5

+

1 User-Based CF ($|N|=2$): Problem

	users											
	1	2	3	4	5	6	7	8	9	10	11	12
1	1		3		?	5			5		4	
2			5	4			4			2	1	3
3	2	4		1	2		3		4	3	5	
4		2	4		5			4			2	
5			4	3	4	2					2	5
6	1		3		3			2			4	



- estimate rating of movie 1 by user 5



User-Based CF ($|N|=2$): Steps

- Who are the similar users ($N=2$) to User5?
 - Normalize matrix by “an average of user rating”
 - Compute cosine similarity

- Which **users**?

- Predict rating of Item1
 - **Average=?**
 - **Weighted Average=?**

$$r_{xi} = \frac{\sum_{j \in N(i;x)} s_{ij} r_{xj}}{\sum_{j \in N(i;x)} s_{ij}}$$

	users											
	1	2	3	4	5	6	7	8	9	10	11	12
1	1		3		?	5			5		4	
2			5	4			4			2	1	3
3	2	4		1	2		3		4	3	5	
4		2	4		5			4			2	
5			4	3	4	2					2	5
6	1		3		3			2			4	

+

2 Item-Based CF ($|N|=2$): Problem

	users											
	1	2	3	4	5	6	7	8	9	10	11	12
1	1			3		?	5			5		4
2				5	4			4			2	1
3	3	2	4		1	2		3		4	3	5
4			2	4		5			4			2
5				4	3	4	2				2	5
6	6	1			3		3		2			4



- estimate rating of movie 1 by user 5

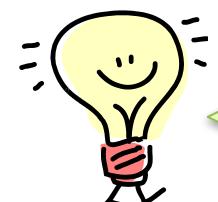


Item-Based CF ($|N|=2$): Steps

- What are the similar items ($N=2$) to Item1?
 - Normalize matrix by “an average of item rating”
 - Compute cosine similarity
 - Which items?

- Predict rating of User5
 - Average=?
 - Weighted Average=?

$$r_{xi} = \frac{\sum_{j \in N(i,x)} s_{ij} r_{xj}}{\sum_{j \in N(i,x)} s_{ij}}$$



	users											
	1	2	3	4	5	6	7	8	9	10	11	12
1	1		3		?	5			5		4	
2			5	4				4		2	1	3
3	2	4		1	2			3	4	3	5	
4		2	4		5			4			2	
5			4	3	4	2					2	5
6	1		3		3			2			4	

- There is still a better prediction solution



Intuitive Idea



- Mean movie rating: **3.7 stars**
- *The Sixth Sense* is **0.5 stars above avg.**
- Joe rates **0.2 stars below** avg.
⇒ **Baseline estimation:**
Joe will rate *The Sixth Sense* 4 stars



		users											
		1	2	3	4	5	6	7	8	9	10	11	12
movies	1	1		3		?	5			5		4	
	2			5	4			4			2	1	3
3	2	4		1	2			3		4	3	5	
4		2	4		5			4			2		
5			4	3	4	2					2	5	
6	1		3	3			2				4		

- What should be **an expected rate** of “User5 on Item3”?
- What should be **deviation from the expectation**?



$$r_{xi} = \frac{\sum_{j \in N(i; x)} s_{ij} r_{xj}}{\sum_{j \in N(i; x)} s_{ij}}$$

CF: Common Practice

- Define **similarity** s_{ij} of items i and j
- Select k nearest neighbors $N(i; x)$
 - Items most similar to i , that were rated by x
- Estimate rating r_{xi} as the weighted average:

$$r_{xi} = b_{xi} + \frac{\sum_{j \in N(i; x)} s_{ij} \cdot (r_{xj} - b_{xj})}{\sum_{j \in N(i; x)} s_{ij}}$$

baseline estimate for r_{xi}

$$b_{xi} = \mu + b_x + b_i$$

- μ = overall mean movie rating
- b_x = rating deviation of user x
 $= (\text{avg. rating of user } x) - \mu$
- b_i = rating deviation of movie i



Item-Based CF ($|N|=2$)

Using Common Approach

- What are the similar items ($N=2$) to Item1?
 - Normalize matrix by “an average of item rating”
 - Compute cosine similarity
 - Which items?

- Predict rating of User5

$$r_{xi} = b_{xi} + \frac{\sum_{j \in N(i;x)} s_{ij} \cdot (r_{xj} - b_{xj})}{\sum_{j \in N(i;x)} s_{ij}}$$

baseline estimate for r_{xi}

$$b_{xi} = \mu + b_x + b_i$$

	users											
	1	2	3	4	5	6	7	8	9	10	11	12
1	1		3		?	5			5		4	
2			5	4			4			2	1	3
3	2	4		1	2		3		4	3	5	
4		2	4		5			4			2	
5			4	3	4	2					2	5
6	1		3		3			2			4	



Remarks

- User-based CF vs. Item-based CF
 - In practice, it has been observed that item-item often works better than user-user
 - Why? Items are simpler, users have multiple tastes

- Major drawbacks in CF
 - Cold Start:
 - Need enough users in the system to find a match
 - Sparsity:
 - Hard to find users that have rated the same items
 - First rater:
 - **New users or items; NO rating data to find neighbors -- cannot recommend!!!**
 - Popularity bias:
 - Cannot recommend items to someone with unique taste
 - Tends to recommend popular items



Evaluation

		movies					
		1	3	4			
			3	5			5
				4	5		5
					3		
					3		
users		2			?		?
						?	
			2	1			?
					?		
		1					

Test Data Set



The diagram illustrates a user-movie rating matrix for evaluation. The columns represent movies and the rows represent users. The matrix contains numerical ratings (1-5) and question marks (?) indicating missing data. A vertical double-headed arrow on the left indicates the rows are labeled 'users'. A horizontal double-headed arrow at the top indicates the columns are labeled 'movies'. An annotation 'Test Data Set' with an arrow points to the bottom-right section of the matrix, which contains all question marks, representing the portion of the data used for testing the recommendation system.



Evaluating Predictions

■ Compare predictions with known ratings

- Root-mean-square error (RMSE)

- $\sqrt{\sum_{xi} (r_{xi} - r_{xi}^*)^2}$ where r_{xi} is predicted, r_{xi}^* is the true rating of x on i

- Precision at top 10:

- % of those in top 10

- Rank Correlation:

- Spearman's correlation between system's and user's complete rankings

■ Another approach: 0/1 model

- Coverage (Recall):

- Number of items/users for which system can make predictions

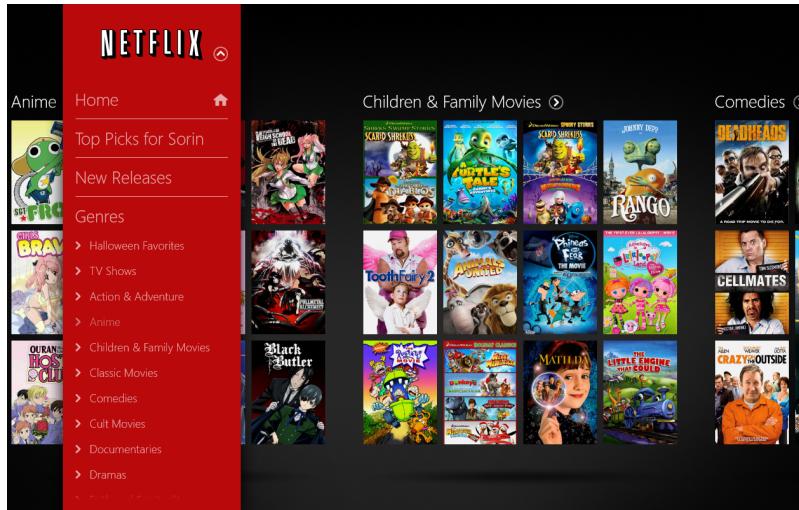
- Precision:

- Accuracy of predictions

- Receiver operating characteristic (ROC)

- Tradeoff curve between false positives and false negatives

- The Netflix Prize
- Further Improvement





The Netflix Prize

■ Training data

- 100 million ratings, 480,000 users, 17,770 movies
- 6 years of data: 2000-2005

■ Test data

- Last few ratings of each user (2.8 million)
- **Evaluation criterion:** Root Mean Square Error (RMSE) = $\frac{1}{|R|} \sqrt{\sum_{(i,x) \in R} (\hat{r}_{xi} - r_{xi})^2}$
- **Netflix's system RMSE: 0.9514**

■ Competition

- 2,700+ teams
- **\$1 million** prize for 10% improvement on Netflix

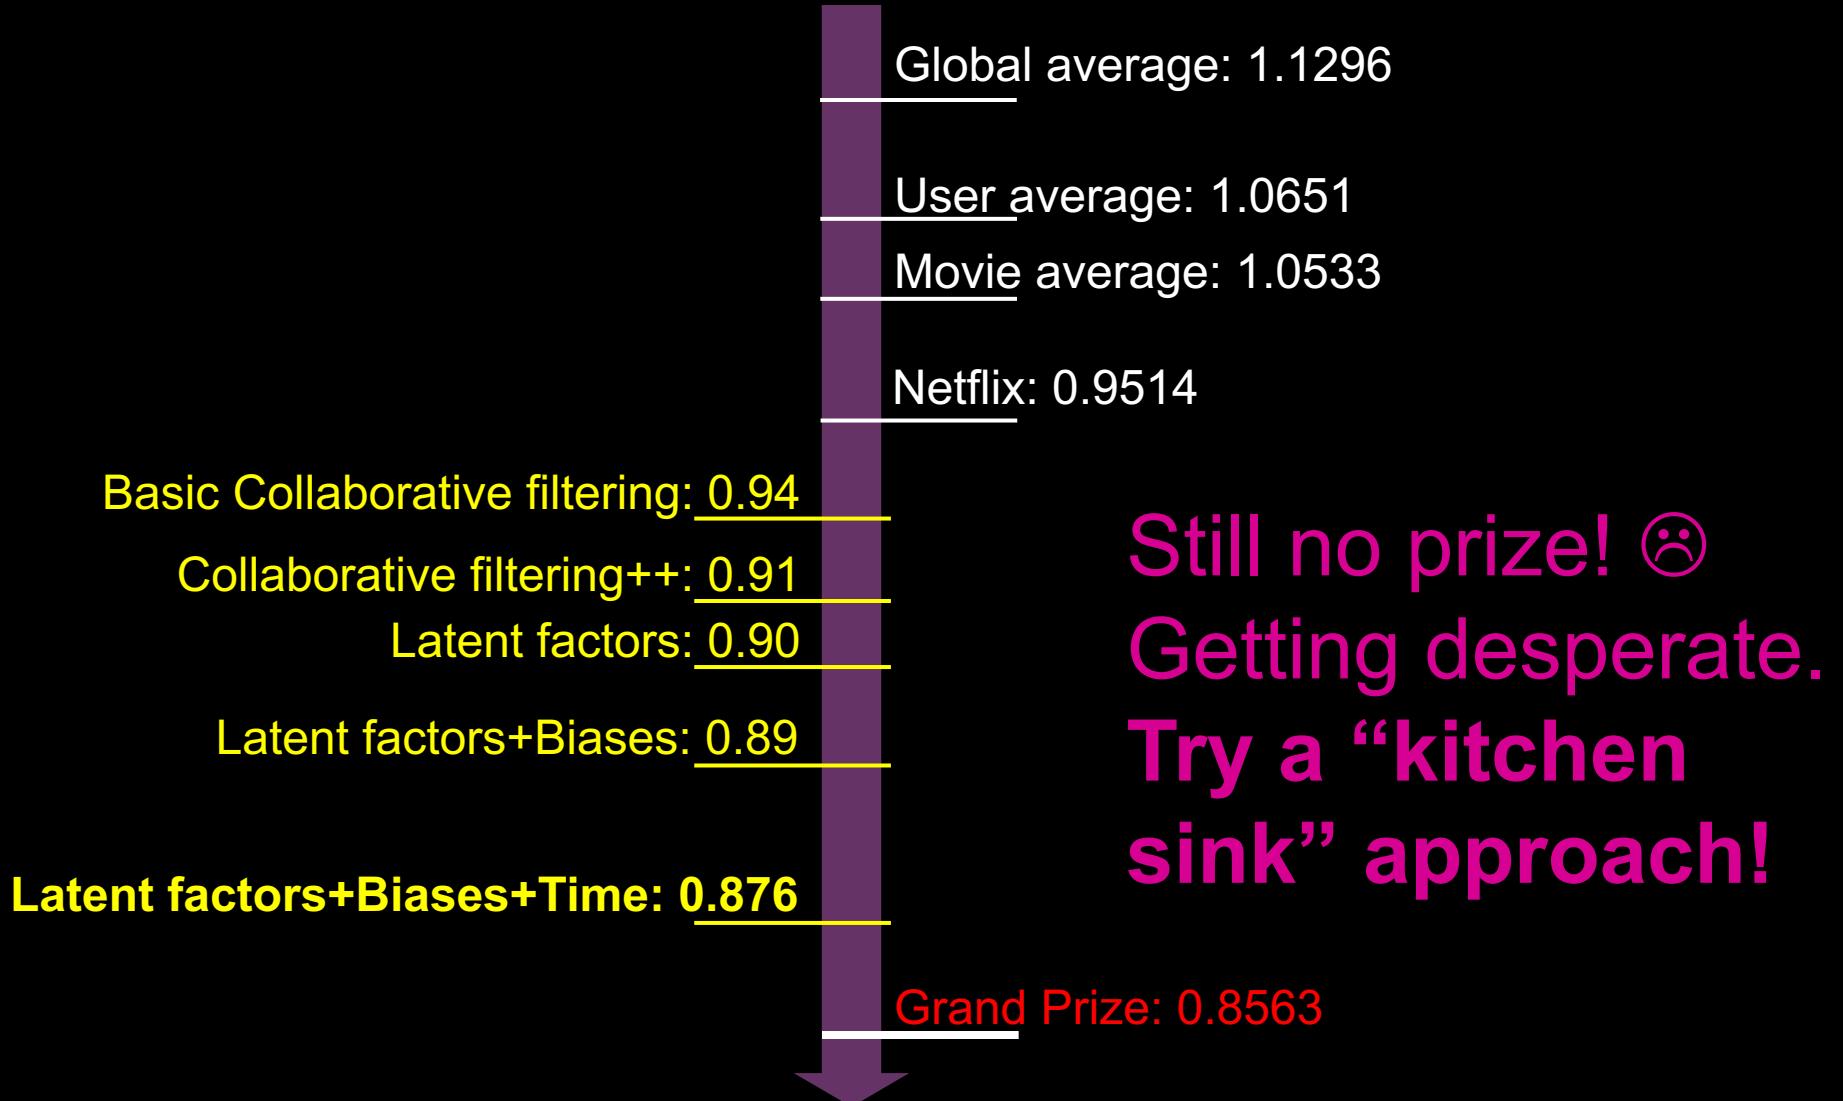
+ The Netflix Utility Matrix R

Matrix R

480,000 users

17,700 movies

1	3	4			
	3	5			5
		4	5		5
			3		
			3		
2			2		2
				5	
	2	1			1
		3		3	
1					



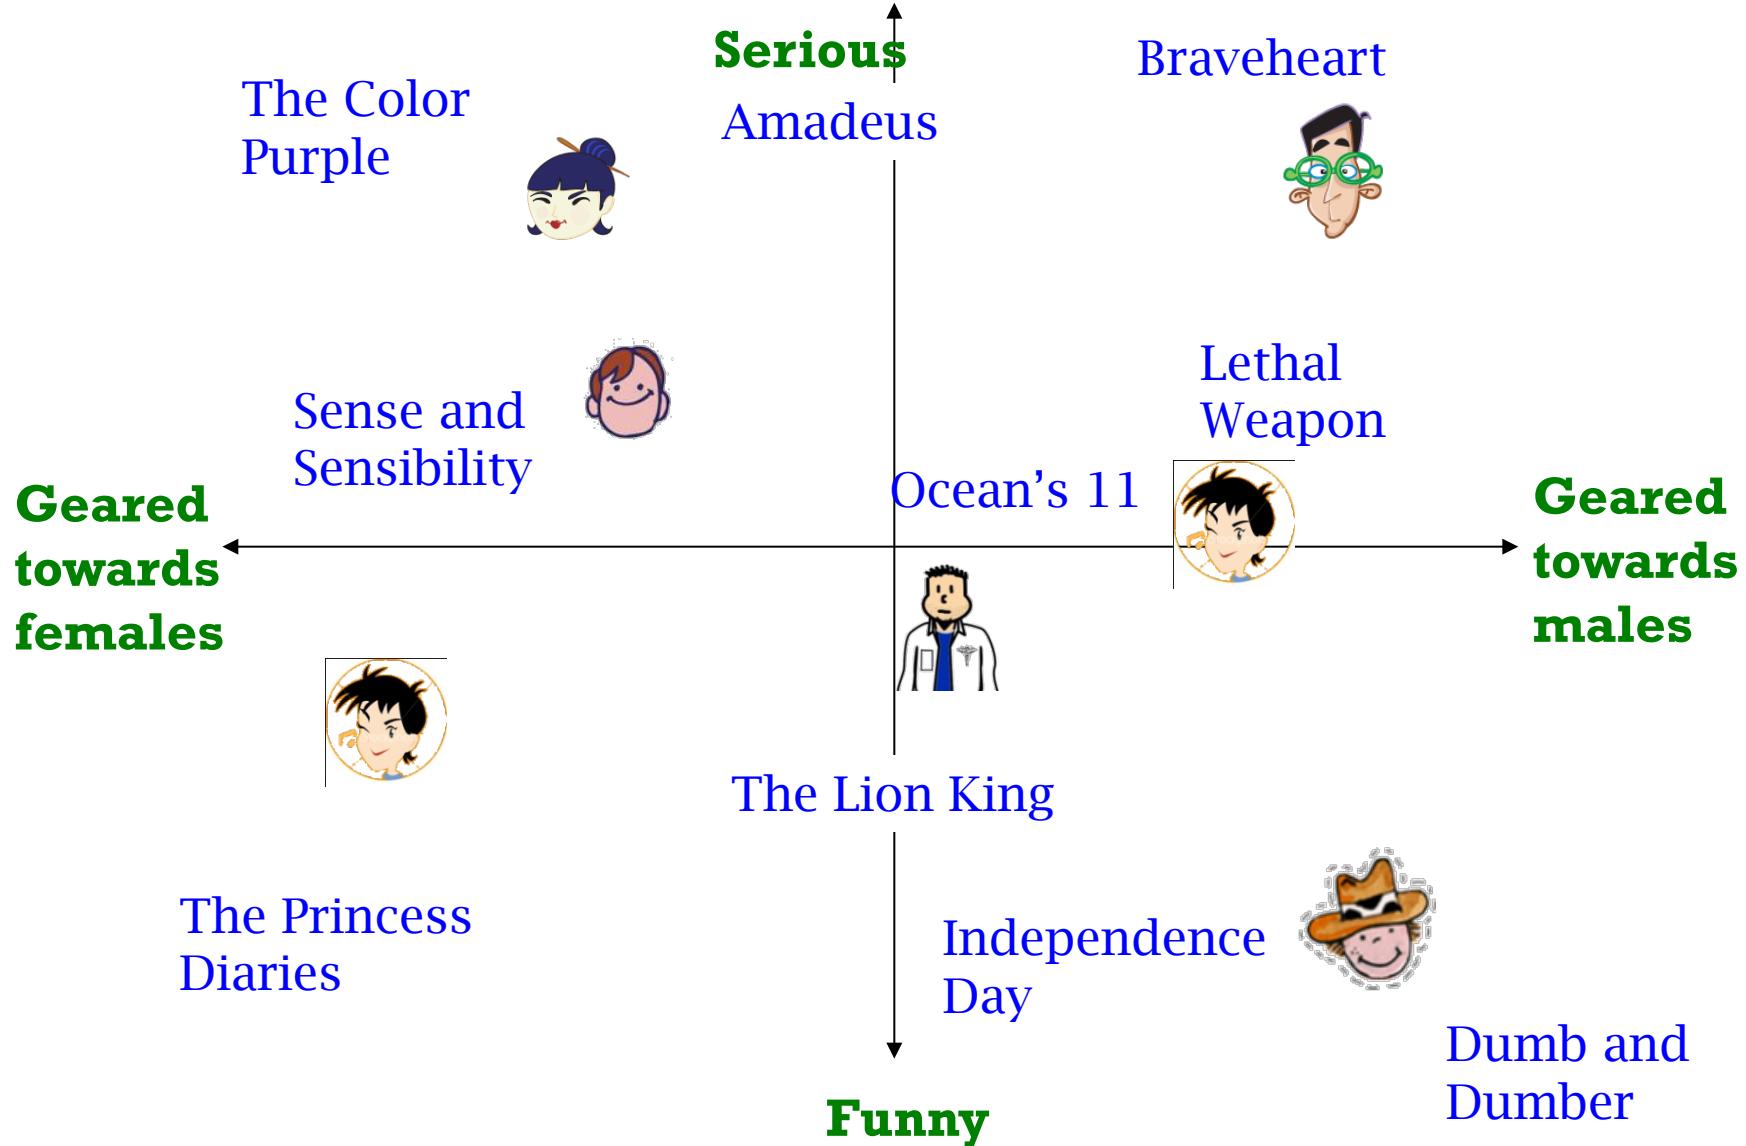


Further Improvement

- Latent Factor
- Temporal Model
- Ensemble



Latent Factor Models (e.g., SVD)



Latent Factor Models

- “SVD” on Netflix data: $\mathbf{R} \approx \mathbf{Q} \cdot \mathbf{P}^T$

$$\begin{matrix} & \text{users} \\ \text{items} & \begin{matrix} 1 & 3 & 5 & 5 & 4 & 2 & 1 & 3 \\ 5 & 4 & 4 & 2 & 3 & 4 & 3 & 5 \\ 2 & 4 & 1 & 2 & 3 & 4 & 3 & 5 \\ 2 & 4 & 5 & 4 & 4 & 2 & 2 & 5 \\ 4 & 3 & 4 & 2 & 2 & 2 & 4 & \\ 1 & 3 & 3 & & 2 & & 4 \end{matrix} \end{matrix} \approx \begin{matrix} & \text{factors} \\ \text{items} & \begin{matrix} .1 & -.4 & .2 \\ -.5 & .6 & .5 \\ -.2 & .3 & .5 \\ 1.1 & 2.1 & .3 \\ -.7 & 2.1 & -2 \\ -1 & .7 & .3 \end{matrix} \end{matrix}$$

R Q

Item (Q) User (P)

$$\text{SVD: } A = U \Sigma V^T$$

users factors

$$\begin{matrix} 1.1 & -.2 & .3 & .5 & -2 & -.5 & .8 & -.4 & .3 & 1.4 & 2.4 \\ -.8 & .7 & .5 & 1.4 & .3 & -1 & 1.4 & 2.9 & -.7 & 1.2 & -.1 \\ 2.1 & -.4 & .6 & 1.7 & 2.4 & .9 & -.3 & .4 & .8 & .7 & -.6 \end{matrix}$$

P^T

Matrix Factorization Techniques for Recommender Systems

Authors: [Yehuda Koren](#) [Yahoo Research](#)
[Robert Bell](#) [AT&T Labs](#)
[Chris Volinsky](#) [AT&T Labs](#)

Published in:

- Journal
[Computer archive](#)
 Volume 42 Issue 8, August 2009
 Pages 30-37
 IEEE Computer Society Press Los Alamitos, CA, USA



2009 Article

- orig-research



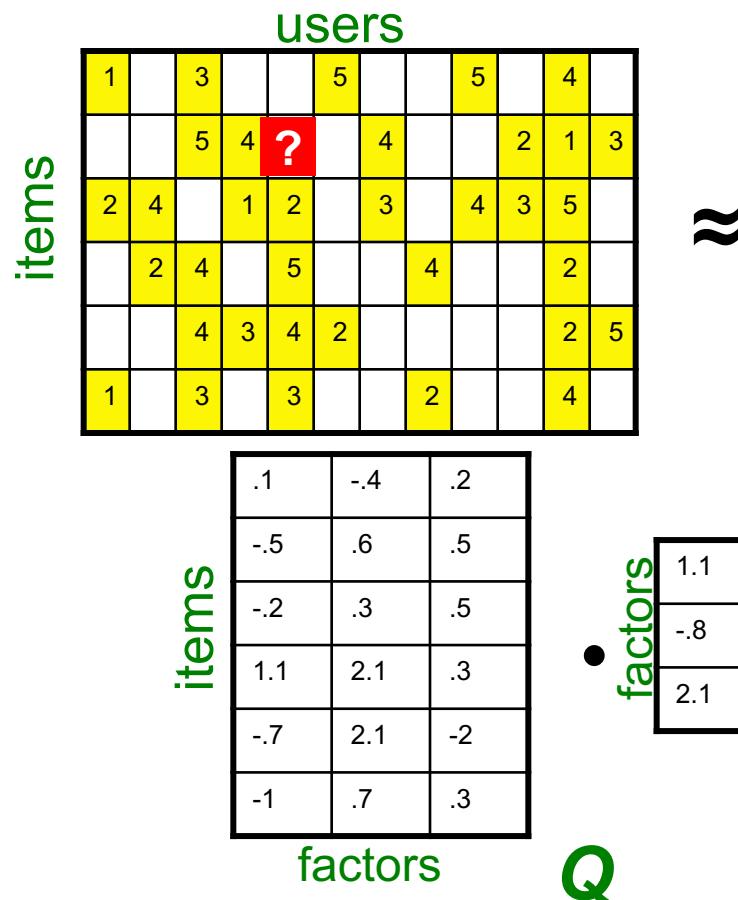
[Bibliometrics](#)

- Downloads (6 Weeks): 0
- Downloads (12 Months): 0
- Downloads (cumulative): 0
- Citation Count: 569



Ratings as Products of Factors

- How to estimate the missing rating of user x for item i ?



$$\begin{aligned}\hat{r}_{xi} &= q_i \cdot p_x \\ &= \sum_f q_{if} \cdot p_{xf}\end{aligned}$$

$q_i = \text{row } i \text{ of } Q$
 $p_x = \text{column } x \text{ of } P^T$

users

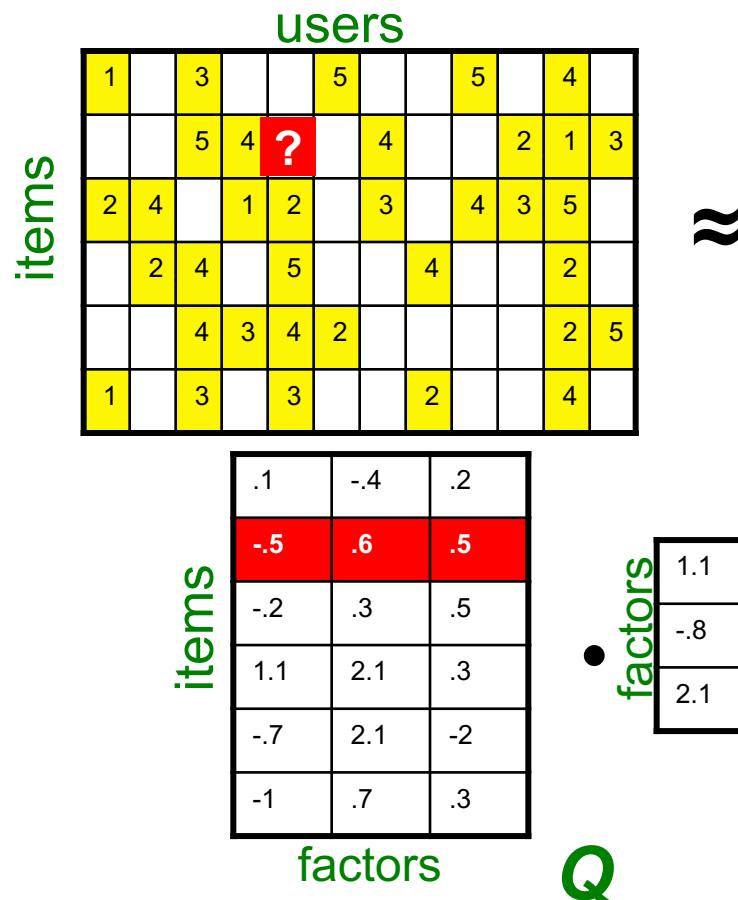
1.1	-.2	.3	.5	-2	-.5	.8	-.4	.3	1.4	2.4	-.9
-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1	1.3
2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1

P^T



Ratings as Products of Factors

- How to estimate the missing rating of user x for item i ?



$$\begin{aligned}\hat{r}_{xi} &= q_i \cdot p_x \\ &= \sum_f q_{if} \cdot p_{xf}\end{aligned}$$

$q_i = \text{row } i \text{ of } \mathbf{Q}$
 $p_x = \text{column } x \text{ of } \mathbf{P}^T$

users

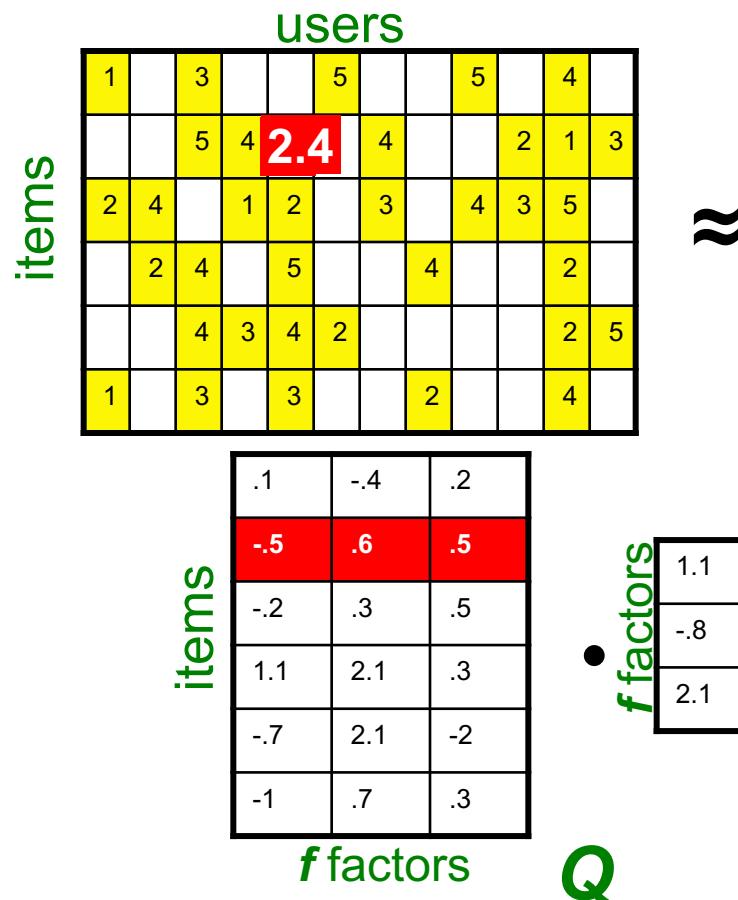
1.1	-.2	.3	.5	-2	-.5	.8	-.4	.3	1.4	2.4	-.9
-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1	1.3
2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1

\mathbf{P}^T



Ratings as Products of Factors

- How to estimate the missing rating of user x for item i ?



$$\begin{aligned}\hat{r}_{xi} &= q_i \cdot p_x \\ &= \sum_f q_{if} \cdot p_{xf}\end{aligned}$$

$q_i = \text{row } i \text{ of } Q$
 $p_x = \text{column } x \text{ of } P^T$

users

1.1	-.2	.3	.5	-2	-.5	.8	-.4	.3	1.4	2.4	-.9
-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1	1.3
2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1

P^T

Temporal Biases & Factors

■ Original model:

$$r_{xi} = \mu + b_x + b_i + q_i \cdot p_x$$

■ Add time dependence to biases:

$$r_{xi} = \mu + b_x(t) + b_i(t) + q_i \cdot p_x$$

- Make parameters b_x and b_i to depend on time
- (1) Parameterize time-dependence by linear trends
(2) Each bin corresponds to 10 consecutive weeks

$$b_i(t) = b_i + b_{i,\text{Bin}(t)}$$

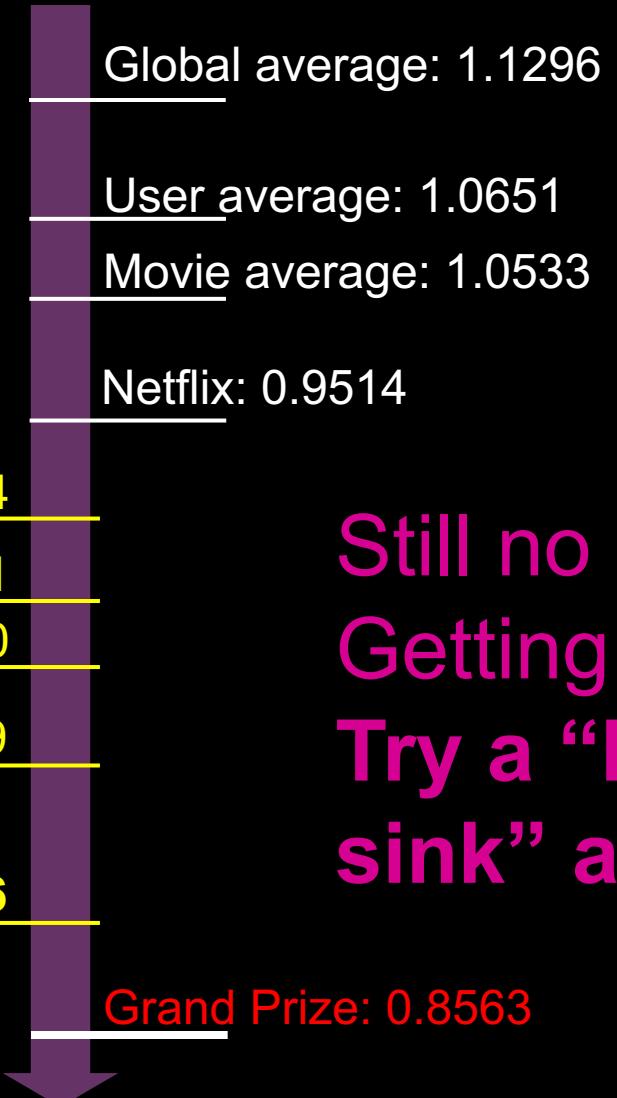
P = user matrix
Q = item matrix
 B_x = user bias
 B_i = item bias

■ Add temporal dependence to factors

- $p_x(t)$... user preference vector on day t

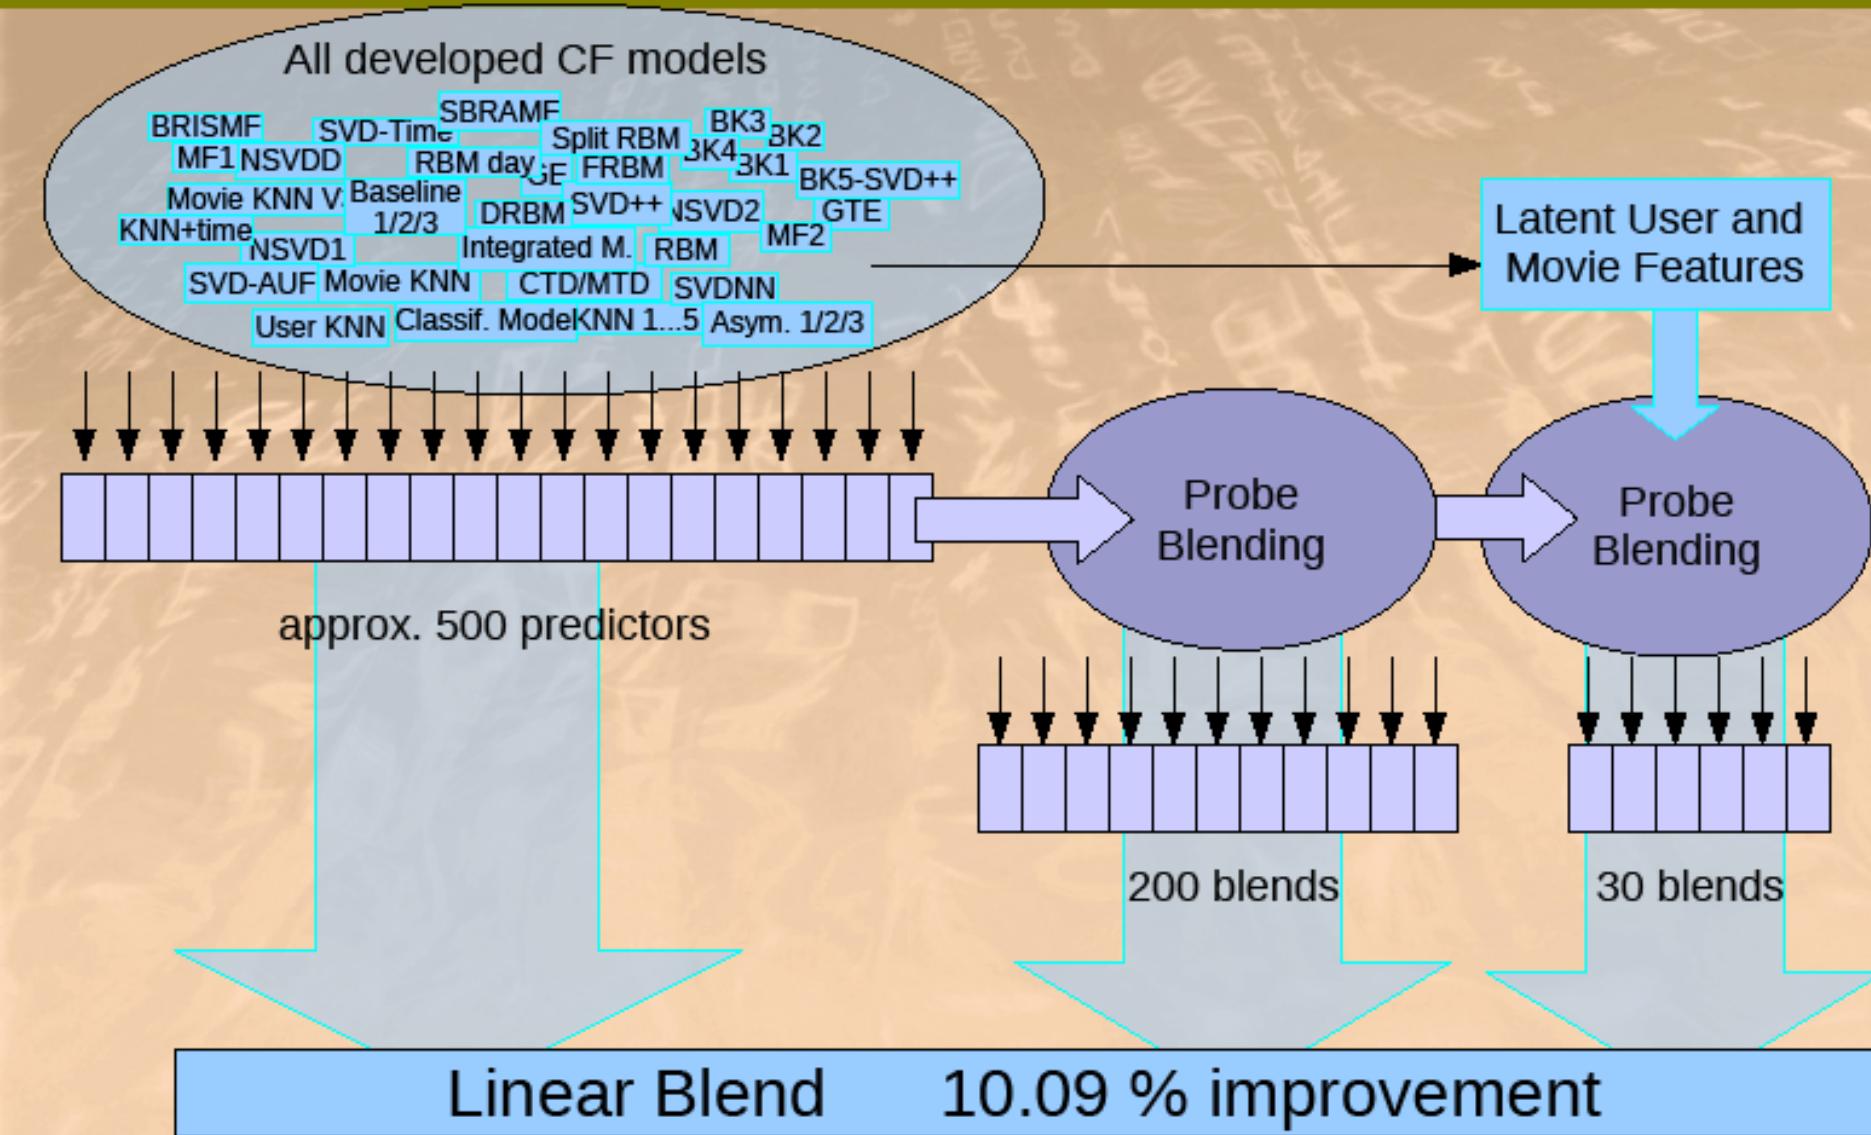
Recap

Basic Collaborative filtering: 0.94
Collaborative filtering++: 0.91
Latent factors: 0.90
Latent factors+Biases: 0.89
Latent factors+Biases+Time: 0.876



Still no prize! 😥
Getting desperate.
Try a “kitchen sink” approach!

Solution of BellKor's Pragmatic Chaos



Netflix Prize

COMPLETED[Home](#) | [Rules](#) | [Leaderboard](#) | [Update](#) | [Download](#)

Leaderboard

Showing Test Score. [Click here to show quiz score](#)Display top leaders.

Rank	Team Name	Best Test Score	% Improvement	Best Submit Time
------	-----------	-----------------	---------------	------------------

Grand Prize - RMSE = 0.8567 - Winning Team: BellKor's Pragmatic Chaos

1	BellKor's Pragmatic Chaos	0.8567	10.06	2009-07-26 18:18:28
2	The Ensemble	0.8567	10.06	2009-07-26 18:38:22
3	Grand Prize Team	0.8582	9.88	2009-07-10 21:24:49
4	Opera Solutions and Vandelay United	0.8588	9.84	2009-07-10 01:12:31
5	Vandelay Industries!	0.8591	9.81	2009-07-10 00:32:20
6	PragmaticTheory	0.8594	9.77	2009-06-24 12:06:56
7	BellKor in BigChaos	0.8601	9.70	2009-05-13 08:14:09
8	Dace	0.8612	9.59	2009-07-24 17:18:43
9	Feeds2	0.8622	9.48	2009-07-12 13:11:51
10	BigChaos	0.8623	9.47	2009-04-07 12:33:59
11	Opera Solutions	0.8623	9.47	2009-07-24 00:34:07
12	BellKor	0.8624	9.46	2009-07-26 17:19:11

Progress Prize 2008 - RMSE = 0.8627 - Winning Team: BellKor in BigChaos

13	xiangliang	0.8642	9.27	2009-07-15 14:53:22
14	Gravity	0.8643	9.26	2009-04-22 18:31:32
15	Ces	0.8651	9.18	2009-06-21 19:24:53
16	Invisible Ideas	0.8653	9.15	2009-07-15 15:53:04
17	Just a guy in a garage	0.8662	9.06	2009-05-24 10:02:54
18	J Dennis Su	0.8666	9.02	2009-03-07 17:16:17
19	Craig Carmichael	0.8666	9.02	2009-07-25 16:00:54
20	acmehill	0.8668	9.00	2009-03-21 16:20:50

Progress Prize 2007 - RMSE = 0.8723 - Winning Team: KorBell

Million \$ Awarded Sept 21st 2009



BellKor's Team

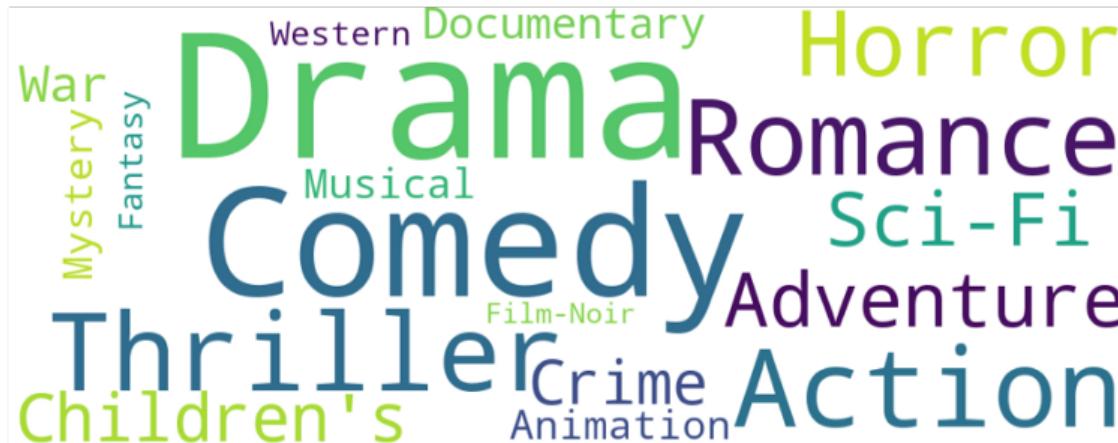


Lab 11: Recommender System

Movie Lens Data Set



- Predict movie rating
 - User-based CF
 - Item-based CF
 - SVD



■ Library

```
import scipy.sparse as sp
from scipy.sparse.linalg import svds

# get SVD components from train matrix. Choose k.
u, s, vt = svds(train_data_matrix, k = 20)
s_diag_matrix=np.diag(s)
X_pred = np.dot(np.dot(u, s_diag_matrix), vt)
print('User-based CF MSE: ' + str(rmse(X_pred, test_data_matrix)))

User-based CF MSE: 2.715056956440157
```

■ Output image

```
In [0]: from sklearn.metrics import mean_squared_error
from math import sqrt
def rmse(prediction, ground_truth):
    prediction = prediction[ground_truth.nonzero()].flatten()
    ground_truth = ground_truth[ground_truth.nonzero()].flatten()
    return sqrt(mean_squared_error(prediction, ground_truth))

In [13]: print('User-based CF RMSE: ' + str(rmse(user_prediction, test_data_matrix)))
print('Item-based CF RMSE: ' + str(rmse(item_prediction, test_data_matrix)))

User-based CF RMSE: 3.1213177574547295
Item-based CF RMSE: 3.4493690282409926
```